

# ECSE 426 – Fall 2015

## Microprocessor Systems

Dept. Electrical and Computer Engineering  
McGill University



McGill

# Problem-oriented Language layer

- Compiled to assembly or instruction set level
- You will be using embedded C
- How does this differ from usual use of C?
  - Directly write to registers to control processor operation
  - All of the registers have been mapped to macros
  - Important bit combinations have macros – use these, please !
  - Registers are 32 bits, so `int` type is 4 bytes
  - Register values may change without your specific instructions
  - Limited output system
  - Floating point operations somewhat inefficient, divide & square-root should be avoided if possible.

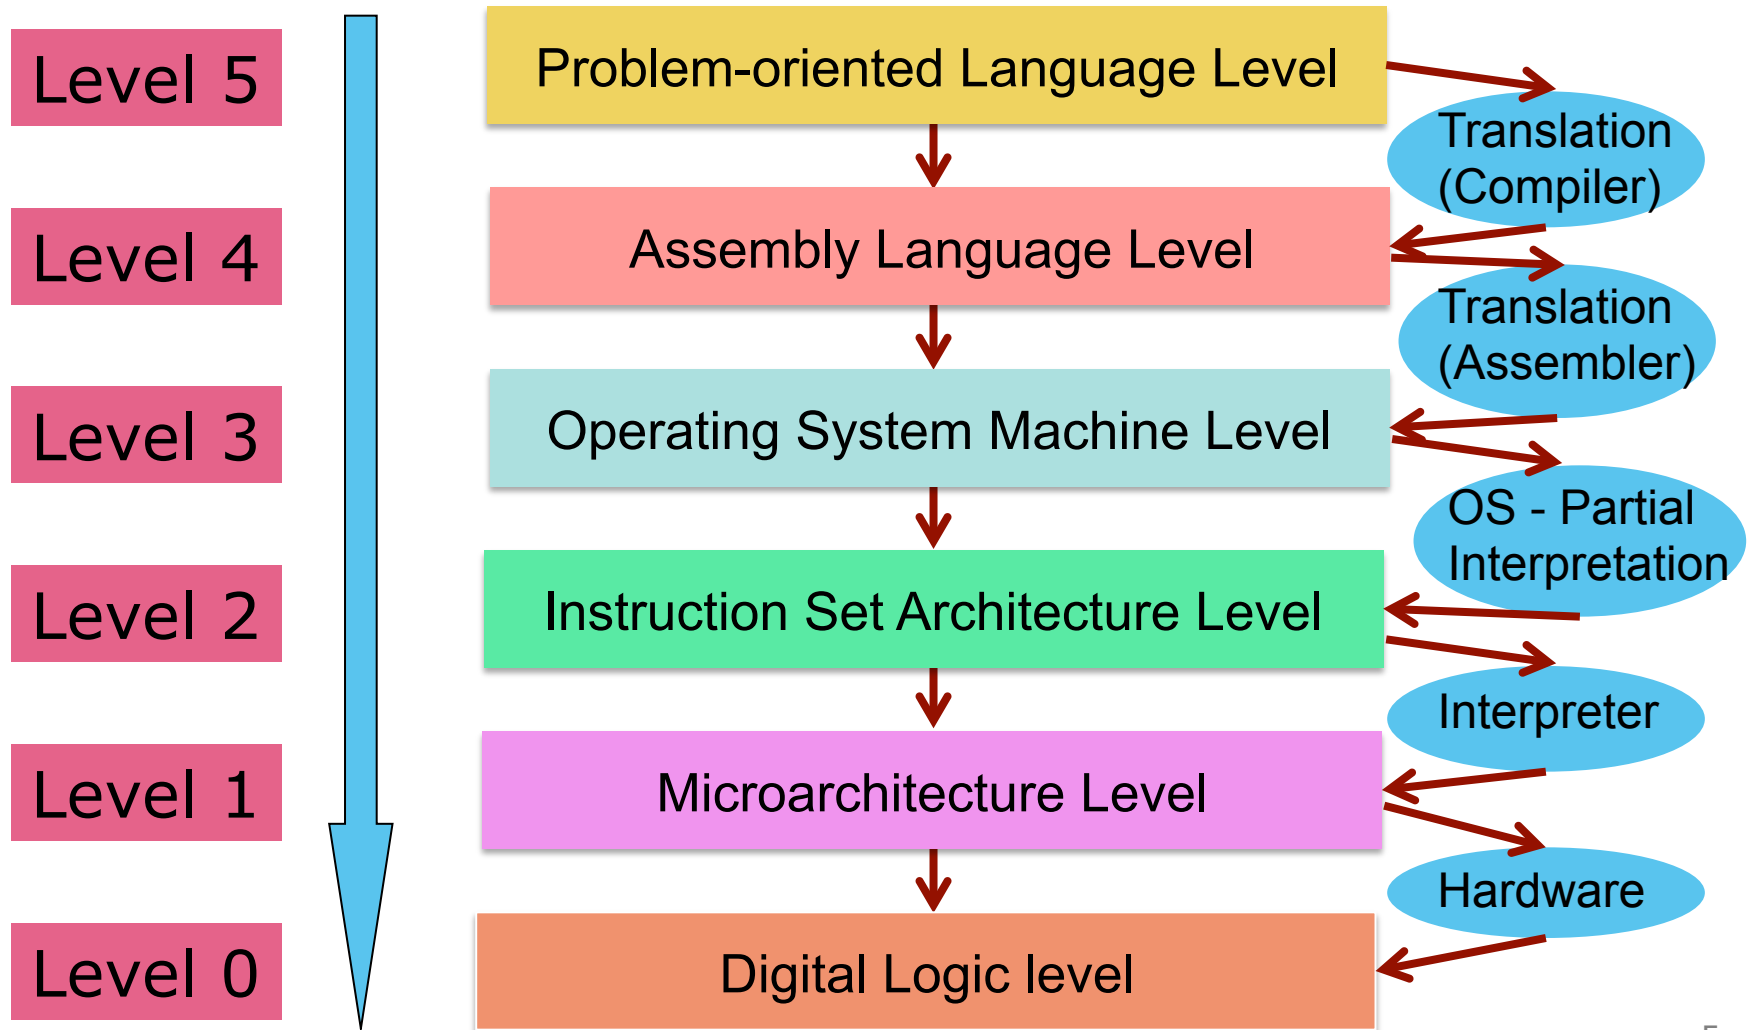
# Assembly versus C

- Efficiency of compiled code
- Source code portability
- Program maintainability
- Typical bug rates (say, per thousand lines of code)
- The amount of time it will take to develop the solution
- Availability and cost of compilers and other development tools
- Your personal experience (or that of the developers on your team) with specific languages or tools
- Don't rule out Java or C++ if you have the memory to play with.

# Problem


- Company “Ostrich” has recently re-developed their embedded software for flagship products
  - Developed in assembly, 80 percent working, 2000 lines of code
  - Suddenly realized that the product isn’t shippable
  - Bugs: system lock-ups indicative of major design flaws or implementation errors & major product performance issues
  - Designer has left the company and provided few notes or comments
- You are hired as a consultant. Do you:
  - Fix existing code?
  - Perform complete software redesign and implementation? In this case, which language?

# Computer Organization



# Instruction Set Architecture

- Interface between HW and SW
  - Virtual Machine
    - Many possible implementations of ISA
- Given by
  - Resources
    - Processor Registers
    - Execution Units
  - Operations
    - Instruction Types
    - Data Types
    - Addressing Modes



i.e., where and how to address operands

# ARM: A Brief History

- Acorn Computers Ltd., Cambridge, England (1983-1985)



- RISC processor concept was introduced in 1980s
- Advanced RISC Machines
- ARM Ltd. – 1990
  - joint venture between Acorn, Apple, and VLSI Technology

# ARM Processor: Architecture

- A 32-bit Enhanced RISC processor
  - RISC: Reduced Instruction Set Computing
    - Simplified instructions can provide higher performance
    - Small, highly-optimized set of instructions
- Register-to-register, three operand instruction set
  - All operands are 32-bits wide
- Employs Load Store Architecture
  - Operations operate on registers and not on memory locations




# Cortex-M3 and Cortex-M4

- 32-bit: internal registers, data path and bus interfaces
- Three-stage pipeline
- Harvard bus architecture with unified memory space
  - instructions and data use same address space
- Based on ARMv7-M architecture
  - high performance processors designed for microcontrollers
  - low power

# ARM Cortex ISA: ARMv7-M

- Implements Thumb2 specification



- Mixture of 32-bit and 16-bit instructions

# ARM Processor Modes


- Processor Mode determines
  - which registers are active (visible)
  - the rights to modify the Program Status Registers (PSRs)
- Each mode can be either **privileged** or **non-privileged**
  - **Privileged** mode
    - full read/write access to PSR
    - The software can use all instructions
  - **Unprivileged** mode
    - Only read access to control fields of PSRs,
    - Read/write access to condition flags
    - Might have restricted access to memory or peripherals

# ARM Modes

- Operation states
  - Thumb state: processor running program code
  - Debug state: halted by debugger; stops executing instructions
- Operation modes
  - Thread mode: default state (privileged to begin with)
  - Handler mode: exceptions and interrupts

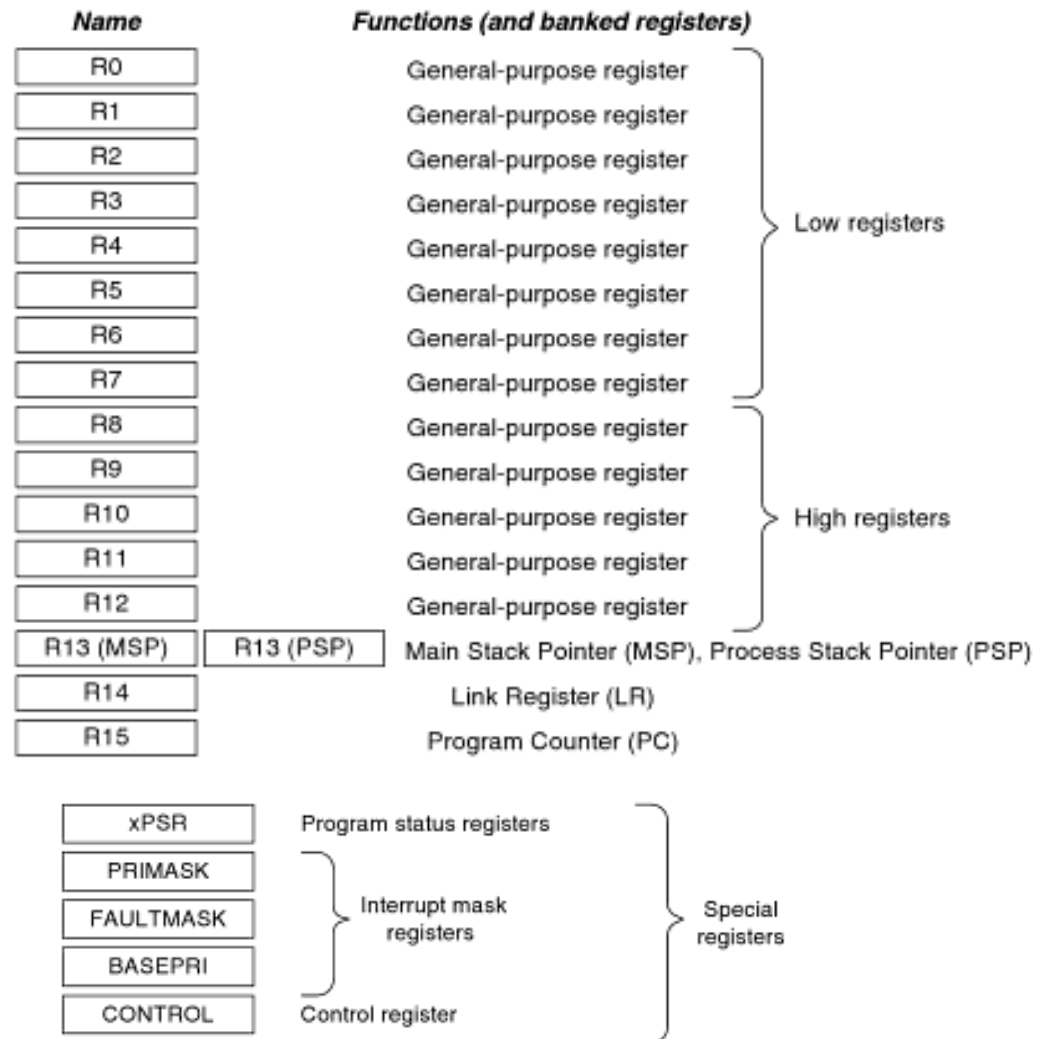
# Registers

- Large register set
  - R0-R12: General purpose
  - Program Counter (R15)
  - Link Register (R14)
  - Stack Pointer (R13)
- Registers R0-R3 hold first 4 words of incoming argument



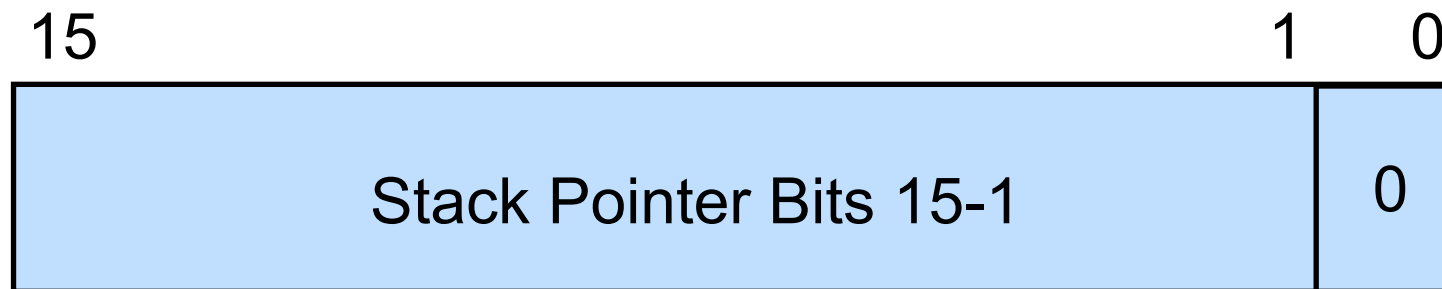
# Registers

- R0-R7
  - low registers
  - Many 16-bit instructions can only access low registers
- R8-R12
  - High registers
  - 32-bit instructions



# R13: Stack Pointer (SP)

- Access stack memory using PUSH and POP operations
- Two different stack pointers
  - Main Stack Pointer (MSP): default, selected after reset
  - Process Stack Pointer (PSP): only in thread mode
- Lowest two bits always zero
- PUSH and POP are 32-bit (addresses must be aligned to 32-bit word boundaries)



# Stack

- PUSH and POP
  - access stack memory
  - Usually used to save (and restore) registers before other data processing
- `PUSH {R0} ; R13 = R13-4, then Memory[R13] = R0`
- `POP {R0} ; R0 = Memory[R13], then R13 = R13 + 4`
- Cortex-M4 uses a full-descending stack arrangement.
- Stack pointer automatically decrements when new data is stored in the stack.



# Stack

- subroutine\_1
  - PUSH {R0-R7, R12, R14} ; Save registers
  - ... ; Do your processing
  - POP {R0-R7, R12, R14} ; Restore registers
  - BX R14 ; Return to calling function
- Can push/pop multiple registers at the same time
- Usually POPs will mirror PUSHes (but there are exceptions)
- PUSH and POP operations are always word aligned
  - addresses must be 0x0, 0x4, 0x8, ...
  - R13 bits 0 and 1 are hardwired to zero and always read as zero

# R14: Link Register (LR)

- Holds return address when calling a function or subroutine
- Program control return to calling program by loading LR into PC
- Function or subroutine call: LR updated automatically
- If a function needs to call another function, needs to save LR on stack
- During exception handling, LR updated automatically to a special EXC\_RETURN value

# R14: Link Register (lr)

```
main ; Main program
```

```
...
```

```
BL function1      ; Call function1 using Branch with Link  
                  ; instruction. PC = function 1 and  
                  ; LR = the next instruction in main
```

```
...
```

```
function1
```


```
...              ; Program code for function 1
```

```
BX LR           ; Return
```

- Bit 0 of the program counter is always 0
- But LR bit 0 is readable and writable
  - Some branch/call instructions require that bit zero is set to one

# R15: Program Counter

- Tells you where you are in the program
  - Next instruction address (pipelining blurs that)
- Readable and writeable
  - read: returns current instruction + 4
  - write: branch operation
- Instructions must be half-word or word-aligned
  - LSB is zero



# R15: Program Counter

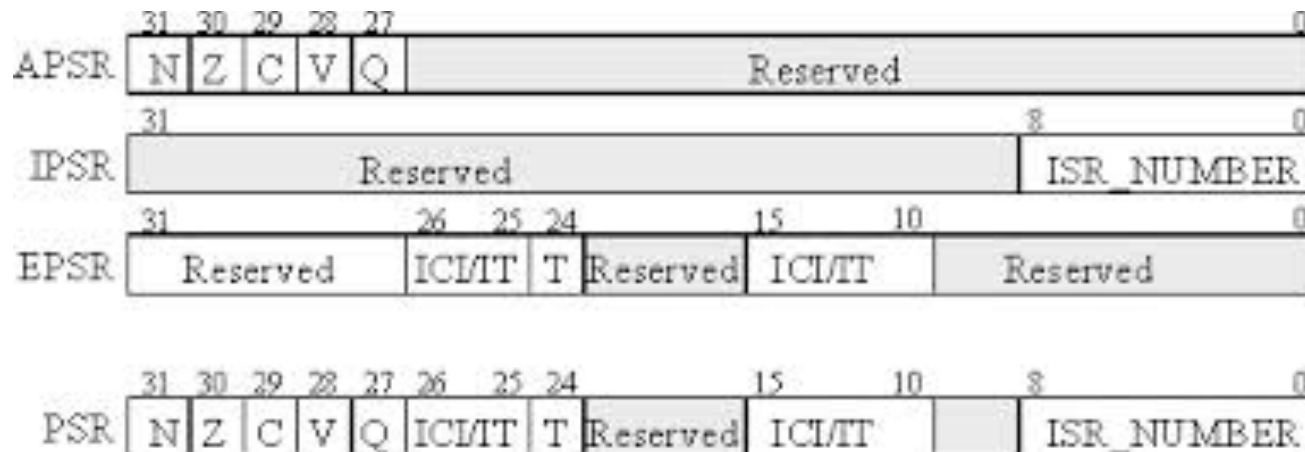
- 32 bits, aligned to even addresses
  - LDR PC, =LABEL      Branch to address in LABEL
  - MOV PC, R4      Branch to address contained in R4
  - MOV R0, PC      Not recommended, processor-dependent
- Need to set LSB to one to indicate Thumb state for some of the branch & read memory instructions
  - handled by compiler automatically in C
- Branches and calls usually have dedicated instructions
- Value of PC useful for accessing literal data stored in program memory

# ARM: Register Summary

- 16 accessible 32-bit registers (R0-R15) at any time
  - R15 acts as the Program Counter (PC)
  - R14 (Link Register) stores subroutine return address
  - R13: stack pointer
- You can (and should!) write in ARM assembly language:
  - PC for R15,
  - lr for R14,
  - sp for R13

# Program Status Register (PSR)

- Three status registers
  - Application PSR (APSR)
  - Execution PSR (EPSR)
  - Interrupt PSR (IPSR)
- Accessed as one combined register (xPSR)



Reproduced from Yiu

# Program Status Register (PSR)

- Access combined PSR
  - `MRS r0, PSR;` read combined program status word
  - `MSR PSR, r0;` write combined program status word
- Or access an individual PSR
  - `MRS r0, APSR;` Read flag state into r0
  - `MRS r0, IPSR;` Read exception/interrupt state
  - `MSR APSR, r0;` write flag state
- EPSR cannot be accessed by code directly using MSR/MRS
- IPSR is read only



# Program Status Register (PSR)

- APSR contains current values of conditional code bits
  - N – Negative Result from ALU
  - Z – Zero result from ALU
  - C – Carry from ALU operation
  - V – ALU operation overflow
- IPSR: Exception number
  - indicates which exception processor is handling
- EPSR:
  - ICI (Interrupt-Continuable Instruction) bits
  - IT: IF-THEN instruction status bit for conditional execution

# Control Register

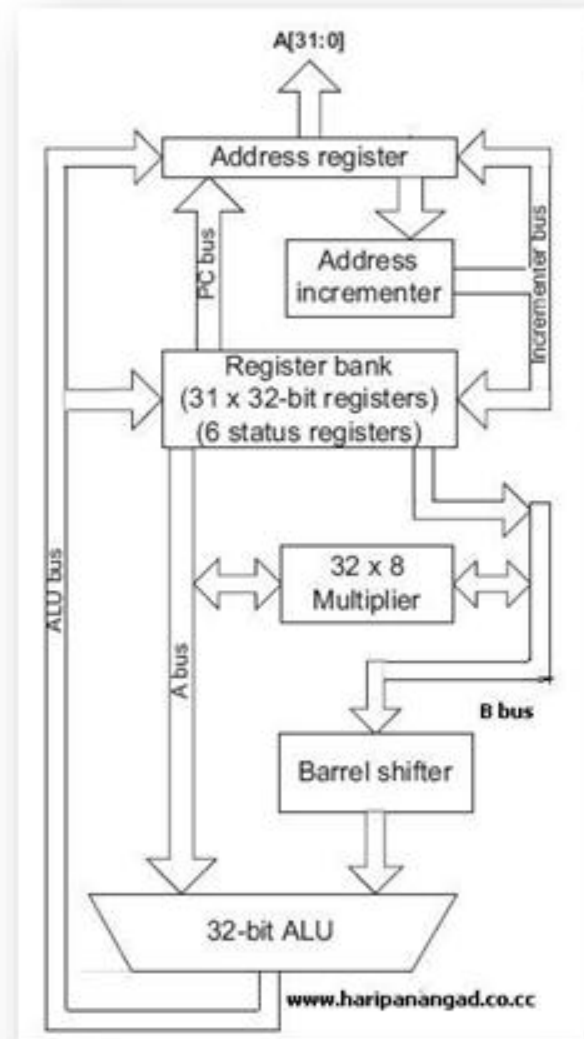
- Access level in thread mode (privileged/unprivileged): bit 0
- Selection of stack pointer (MSP/PSP): bit 1
- Current context uses floating point unit: bit 2
- All zeros: privileged, MSP, FPU unused
- Can only be modified in privileged access level
- Switch back to privileged access through an exception
  - exception handler can clear the nPRIV bit

# Floating Point Registers

- Cortex M4 has additional registers for floating point operations
- S0:S31 can be accessed using floating point instructions
- or accessed as pairs (D0:D15)
- Does not support double precision floating point ops but can transfer double precision data


# ARM: Core Datapath

- Separate control and datapath
- In datapath operands are not fetched directly from memory locations
  - Data is placed in register files
  - No data processing takes place in memory
- Instructions typically use 3 registers
  - 2 source registers
  - 1 destination register
- Barrel shifter preprocesses data before it enters ALU



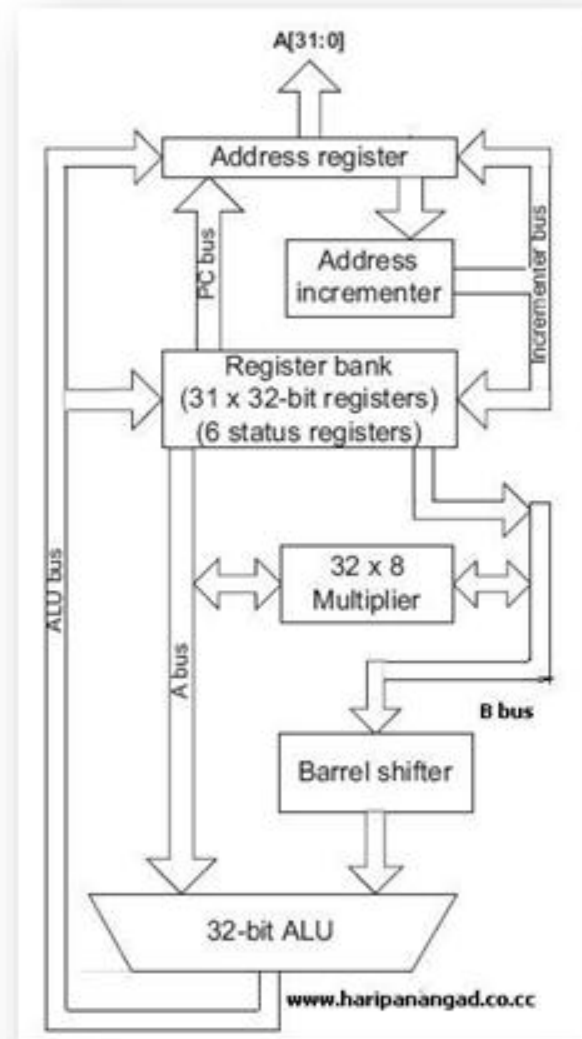
# ARM Organization

- Data has to be moved from the memory location to the central set of registers
  - After processing it is stored back in memory
- Register bank connected to ALU via two datapath buses A and B
  - Bus B goes through Barrel shifter
    - It pre-processes data from source register by shl, shr or rotation
  - PC is part of the register bank
    - responsible for generating addresses for next operation
  - Registers can handle both data and address



# ARM Organization

- Address Incrementer
  - block increments or decrements register values independent of ALU



# Cortex-M feature set comparison

	Cortex-M0	Cortex-M3	Cortex-M4
Architecture Version	V6M	v7M	v7ME
Instruction set architecture	Thumb, Thumb-2 System Instructions	Thumb + Thumb-2	Thumb + Thumb-2, DSP, SIMD, FP
DMIPS/MHz	0.9	1.25	1.25
Bus interfaces	1	3	3
Integrated NVIC	Yes	Yes	Yes
Number interrupts	1-32 + NMI	1-240 + NMI	1-240 + NMI
Interrupt priorities	4	8-256	8-256
Breakpoints, Watchpoints	4/2/0, 2/1/0	8/4/0, 2/1/0	8/4/0, 2/1/0
Memory Protection Unit (MPU)	No	Yes (Option)	Yes (Option)
Integrated trace option (ETM)	No	Yes (Option)	Yes (Option)
Fault Robust Interface	No	Yes (Option)	No
Single Cycle Multiply	Yes (Option)	Yes	Yes
Hardware Divide	No	Yes	Yes
WIC Support	Yes	Yes	Yes
Bit banding support	No	Yes	Yes
Single cycle DSP/SIMD	No	No	Yes
Floating point hardware	No	No	Yes
Bus protocol	AHB Lite	AHB Lite, APB	AHB Lite, APB
CMSIS Support	Yes	Yes	Yes

# ARM Cortex-M3/4 Processor

- Complex processor core providing standard CPU and system architecture
- Architecture v7-M
  - Only 2 processor modes (thread/handler)
  - Processor automatically saves/restores state in exceptions
    - Improved interrupt speeds
  - Enhanced Flash access



# ARM Cortex-M3/4 Processor

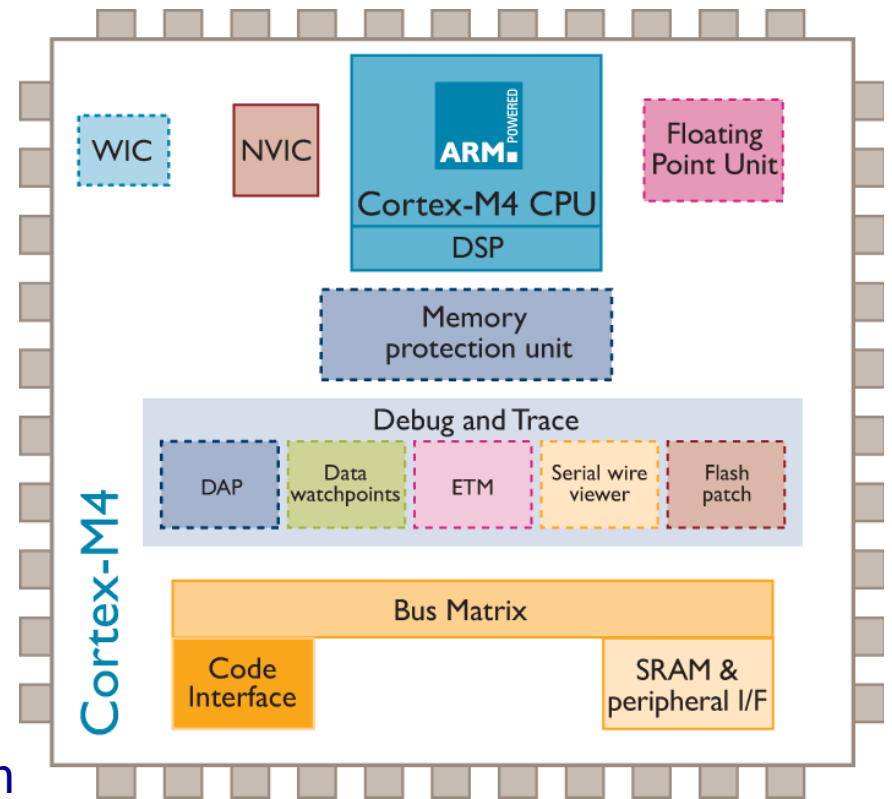
- Banked Stack Pointer (SP) (banked: only one visible at a time)
  - Main Stack Pointer (MSP): The default stack pointer, used by the operating system (OS) kernel and exception handlers
  - Process Stack Pointer (PSP): Used by user application code
  - Stack usage by OS kernel can be separated from application usage: better reliability and optimized stack space usage
  - No OS: Can use the MSP all the time

# Cortex-M

- Basic peripherals are built into Cortex-M, e.g.
  - interrupt controller
  - 24-bit SysTick timer
  - debug system
  - memory map and sleep mode
- Based on **Harvard architecture** (code and data bus separated)
  - Enables simultaneous instruction fetch with data load/store
- **Multiple buses** for performing operations in **parallel**
- **Bus matrix** connects processor and debug interface to external busses

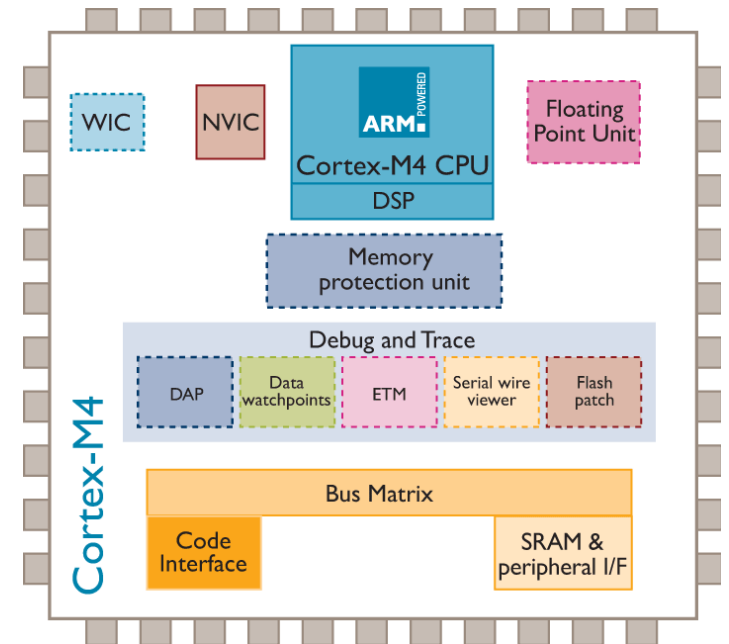
# Cortex-M4

- ARMv7ME Architecture
  - Thumb-2 Technology
  - Compatible with Cortex-M3
  - Single precision FPU
  - DSP and SIMD extensions
  - Single cycle MAC ( $32 \times 32 + 64 \rightarrow 64$ )
- Microarchitecture
  - 3-stage pipeline with branch speculation
  - 3x AHB-Lite Bus Interfaces




# Cortex-M4

- Configurable for ultra low power
  - Deep Sleep Mode, Wakeup Interrupt Controller
  - Power down features for Floating Point Unit
- Flexible configurations for wider applicability
  - Configurable Interrupt Controller (1-240 Interrupts and Priorities)
  - Debug & Trace
  - Optional Memory Protection Unit



# Cortex-M4

- CPU core is closely coupled to the interrupt controller (NVIC) and various debug logic blocks
- CM4Core: The Cortex-M4 core contains the registers, ALU, data path, and bus interface.
- Nested Vectored Interrupt Controller (NVIC)
  - built-in interrupt controller.
  - Number of interrupts is customized by chip manufacturers (up to 240)
- Bus interconnect

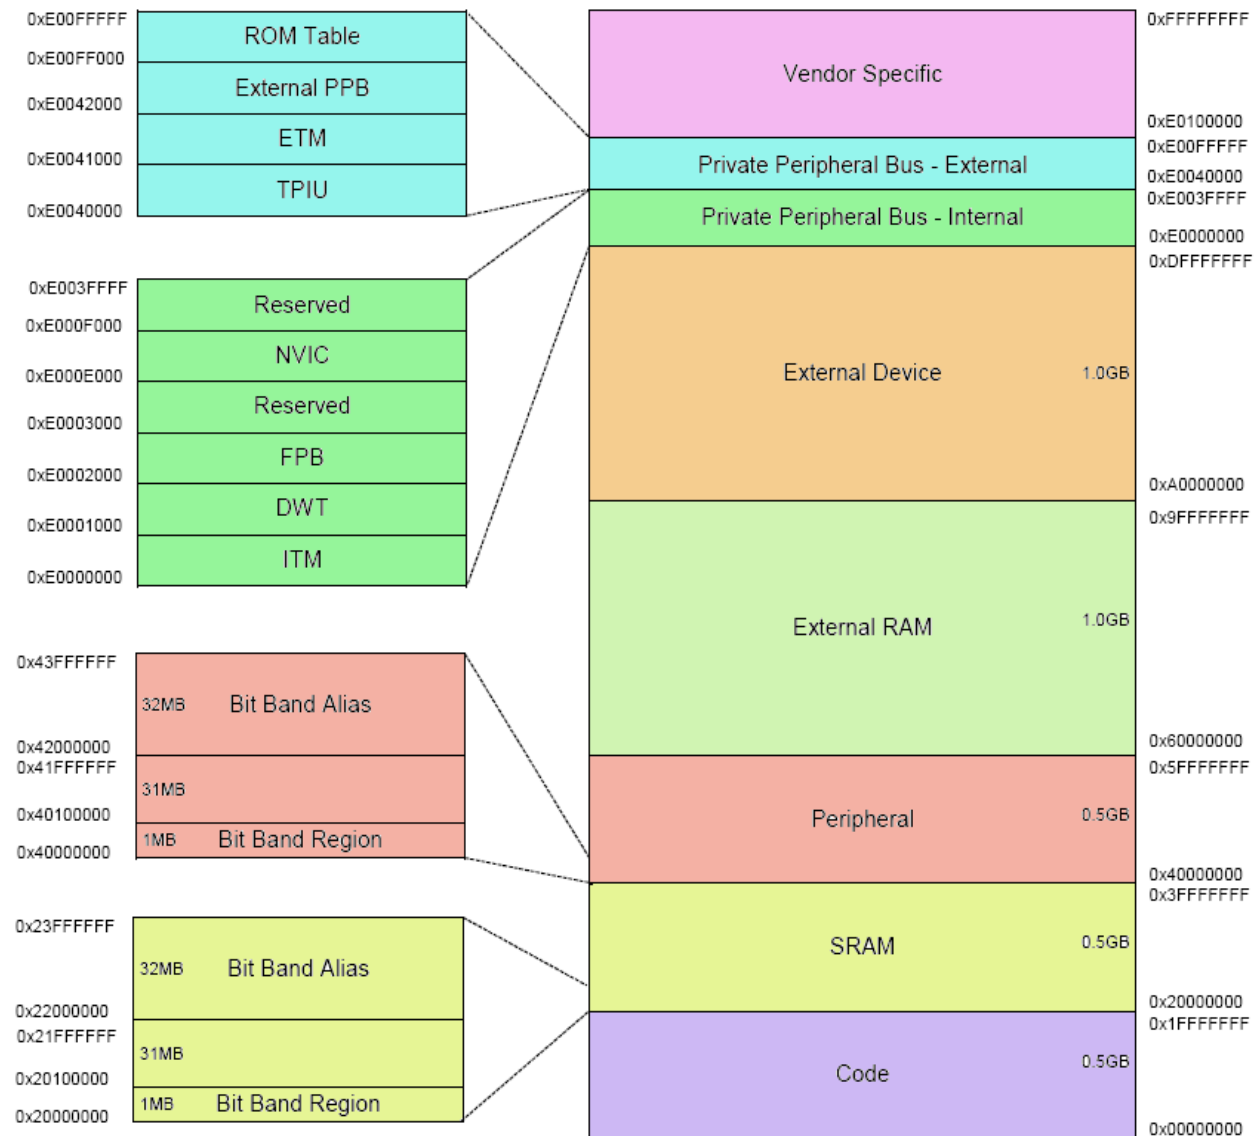


# Cortex M Address Space

- 4 GB address space
  - (ROM/Flash for Code)
- Built-in and external RAM
- Built-in and external IO
  - Memory-mapped IO
  - Special Function Registers (SFRs) accessed by processor

0xFFFFFFFF	System level	Private peripherals including build-in interrupt controller (NVIC), MPU control registers, and debug components
0xE0000000		
0xDFFFFFFF	External device	Mainly used as external peripherals
0xA0000000		
0x9FFFFFFF	External RAM	Mainly used as external memory
0x60000000		
0x5FFFFFFF	Peripherals	Mainly used as peripherals
0x40000000		
0x3FFFFFFF	SRAM	Mainly used as static RAM
0x20000000		
0x1FFFFFFF	CODE	Mainly used for program code. Also provides exception vector table after power up
0x00000000		

# Memory Map



# Cortex Memory Regions

- Program code located in **code**, SRAM or external RAM regions
  - Best to put program code in code region
    - instruction fetches and data accesses are carried out simultaneously on two separate bus interfaces
- **SRAM memory (0x20000000 – 0x3FFFFFFF)** : connecting internal SRAM
  - Executable: program can be copied and executed from this part of memory
  - Access through system interface bus
- **On-chip peripherals (0x40000000 – 0x5FFFFFFF)** is a 0.5 GB block intended for peripherals
  - Code cannot be executed from this region




# Cortex Memory Regions

- Two slots of 1 GB memory space allocated for external RAM and devices
  - External RAM regions
    - intended for either on-chip or off-chip memory
    - Code can be executed in this region
  - External devices
    - intended for external devices and/or shared memory.
    - Non-executable region
- System region: 0.5GB memory (0xE0000000 – 0xFFFFFFFF)
  - System-level components
  - Internal peripheral buses & external peripheral bus
  - Vendor-specific system peripherals
  - Non-executable region

# Instruction Set Architecture

- Machine code vs. assembly
  - Machine code is a string of bits that processors “read”.
    - practically unreadable for humans
  - Assembly is intended for human reading
- In assembly the following instruction formatting is common:

```
label opcode operand1, operand2, ... comments
```



# Assembly Language

```
label opcode operand1, operand2, ... comments
```

- **Label** is optional
- **Opcode** is the actual instruction
- Depending on the instruction there can be two or more **operands**
  - Usually, the first operand is the destination of the operation

# Reset

- After the processor exits reset, it will read two words from memory:
  - Address 0x00000000: Starting value of R13 (the stack pointer)
  - Address 0x00000004: Reset vector (starting address of program execution)
- Initial value for the MSP is put at the beginning of the memory map
- Followed by the vector table (contains vector address values)
- The first item in the vector table (exception type 1) is the reset vector
  - Second piece of data fetched by the processor after reset.


# Moving Data

- Basic function: transfer of data
  - Register to register
  - Memory to register
  - Special register to register
  - Moving an immediate data value into a register
- The command to move data between registers is MOV (move)
  - `MOV R8, R3`
  - `MVN R8, R3 ; (move negative of value in R8)`
- Basic instructions for accessing memory are Load and Store
  - LDR: from memory to registers; STR: from registers to memory
  - Transfers can be in different data sizes (byte, half word, word, double word)

# Moving Data

- Multiple Load/Store operations can be combined into single instructions
  - LDM (Load Multiple) and STM (Store Multiple)
  - e.g. Store multiple words to memory location specified by *Rd*
  - `STMIA.W R8!, {R0-R3} ; R8=0x8000 changed to 0x8010 after`
  - `STMIA.W R8, {R0-R3} ; store (increment by 4 words)`
  - `STMIA.W R8, {R0-R3} ; R8 unchanged after store`
- Memory accesses with pre-indexing and post-indexing.
  - Pre-indexing : register adjusted & updated address is used
  - `LDR.W R0,[R1, #offset]! ; Read memory[R1+offset], with`  
`; R1 updated to R1+offset`
  - Post-indexing : register adjusted after address is used
  - `LDR.W R0,[R1], #offset ; Read memory[R1], with R1`  
`; updated to R1+offset`

# Moving Data

- Moving immediate data into a register is common
  - 8 bits: `MOV R0, #0x12 ; Set R0 to 0x12`
  - Larger values (Thumb 2): `MOVW.W R0, #0x789A ; Set R0 to 0x789A`
  - 32-bit can use two instructions to set the upper and lower halves:
    - `MOVW.W R0, #0x789A ; Set R0 lower half to 0x789A` 
    - `MOVT.W R0, #0x3456 ; Set R0 upper half to 0x3456. Now  
; R0=0x3456789A`

# Store Timing

```
STR Rd, [Rn, #offset]
```

```
Effect: Rd <- [[Rn]+offset]
```

- Execution takes always one cycle
  - Address generation is performed in the initial cycle
  - Data store is performed at the same time as next instruction executes



# Data Processing

- Many data operation instructions can have multiple instruction formats.
  - `ADD R0, R1 ; R0 = R0+R1`
  - `ADD R0, #0x12 ; R0 = R0 + 0x12`
  - `ADD.W R0, R1, R2 ; R0 = R1+R2`

# Arithmetic & Conditional Execution

- Mostly 3-operand (two inputs and output)

XYZ Rd, Rn, <Op2>

- Example:
  - ADD – add
  - ADC – add with carry
  - SUB – subtract, etc.

Examples:

ADD R0, R1, R2;

R0 = R1+R2, no flag change

ADDS R0, R1, R2;

as above, but CNZV affected

ADDCSS R0, R1, R2;

if C flag then ADDS

# Branch Instructions

- Branch can be conditional or unconditional
- Syntax

```
B [cond] label
```

- Example
  - `B label ; Branch to a labeled address`
  - `BX reg ; Branch to address specified by a register`
- Calling a function
  - `BL label ; Branch to a labeled address and save  
; return address in LR`
  - `BLX reg ; Branch to an address specified by a  
; register and save return address in LR`
  - Function can be terminated using `BX LR`
    - program control then returns to the calling process

# Branch Instructions

- BL instruction destroys the current content of your LR register.
- If program code needs the LR register later, LR must be saved
- The common method is to push the LR to stack

```
main
    ...
    BL functionA
    ...
functionA
    PUSH {LR} ; Save LR content to
                ; stack
    ...
    BL functionB
    ...
    POP {PC} ; Use stacked LR content
                ; to return to main
functionB
    PUSH {LR}
    ...
    POP {PC} ; Use stacked LR content
                ; to return to functionA
```

# Conditional Branch Instructions

- Flags in the Current Program Status Register (CPSR) determine whether a branch should be carried out.
- Five flag bits: Four used for branch decisions

Flag	PSR bit	Description
N	31	Negative flag (last operation result is a negative value)
Z	30	Zero (last operation result returns a zero value)
C	29	Carry (last operation returns a carry out or borrow)
V	28	Overflow (last operation results in an overflow)

# Conditional Execution Example

```
int gcd(int a, int b){  
    while (a != b)  
        if(a>b) a= a-b  
        else b = b-a;  
    return a;}  

```


```
gcd    CMP    R0, R1;  
        BEQ    end  
        BLT    less  
        SUB    R0, R0, R1  
        B      gcd  
less   SUB    R1, R1, R0  
        B      gcd  
end    ...
```

```
gcd    CMP    R0, R1;  
        SUBGT   R0, R0, R1; SUB if R1 > R0  
        SUBLT   R1, R1, R0; SUB if R1 < R0  
        BNE     gcd;
```

# ARM Instruction Set - Summary

- All instructions are 32 bit long, many execute in one cycle
- Instructions can be conditionally executed
- Example: data processing instruction

```
if (R0 equal R1) then {  
    R3 = R4 + R5  
    R3 = R3 / 2  
}  
else {  
    R3 = R6 + R7  
    R3 = R3 / 2  
}
```

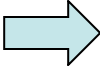
```
CMP R0, R1 ; Compare R0 and R1  
ITTEE EQ ; If R0 equal R1, Then-Then-Else-Else  
ADDEQ R3, R4, R5 ; Add if equal   
ASREQ R3, R3, #1 ; Arithmetic shift right if equal  
ADDNE R3, R6, R7 ; Add if not equal  
ASRNE R3, R3, #1 ; Arithmetic shift right if not equal
```

# Thumb Instruction Set

- Thumb is a 16-bit instruction set
  - Optimized for code density from C code (~65% of ARM code size)
  - Improved performance from narrow memory
  - Subset of functionality of ARM instruction set
- Thumb is not a “regular” instruction set
  - Constraints are not generally consistent
  - Targeted at compiler generation, not hard coding



# Assembler

- **Utility program**: translate **assembly language** into **machine code**
  - Translation is nearly **isomorphic** (one-to-one)
  - Assembly mnemonic statements  machine instructions and data
- Translation of a single high-level language instruction into machine code generally results in many machine instructions
- Assembler may provide **pseudo-instructions**
  - expand into several machine language instructions
  - Example:
    - Assembly language does not support “branch if greater or equal” instruction
    - Assembler can provide pseudo-instruction that expands to “set if less than” and “branch if zero”

# ARM Assembly Language

- Assemblers: nearly as powerful as high-level languages
- Assembly code: **commands** or **directives** (commands to the assembler tool rather than the processor)

Format:

```
{label}{instruction | directive | pseudo-instruction} {;  
    comment}
```

Notes:

- Everything is optional
- Label must start the line (no spaces or tabs before)
- Spaces and tabs freely used otherwise
- Comments can't spread over multiple lines

# Example: Assembly Directives, cont.

- **AREA** directive instructs assembler to assemble a new code/data section
- **Section**: independent, named, indivisible chunks of code (manipulated by linker)
- Syntax `AREA sectionname { , attr}{,attr}...`
  - sectionname: name given to the section. Any can be chosen
- **CODE**: keyword indicating that the section to follow is code
- **EXPORT**: keyword indicating functions, variables, etc. visible to external segment
- **ENTRY**: keyword indicating the entry point of a segment

# Example: Assembly Directives

```
AREA    PROGRAM, CODE
; takes 2 7-digit BCD numbers and places results in the first
; bcdadd(a, b)-> a' with a'= a+b for encoding below
; svxxD6 D5D4 D3D2 D1D0 where s=1 is negative, v=1 is
; overflow, and 7 4-bit BCD Di
; ARM Calling convention: arg1 and arg2 in R0 and R1,
; respectively. Result is in R0 & other registers unchanged.
; no memory used; return via LR register
; position-independent code, needs to be linked with C routine
EXPORT bcdadd
ENTRY
bcdadd ; labels entry point to subroutine & rest of code after
...
END
```

# Assembler Directives

- Important directives for data definition and storage
  - Symbolic constant (to be replaced throughout code): EQU
    - EQU directive gives symbolic name to:
      - numeric constant, a register-relative value or a PC-relative value
  - Variable: DCD, DCB
    - allocate one or more words of memory
    - defines the initial runtime contents of the memory
    - DCD: aligned on four-byte boundaries
  - Register variable: RN (define a name for a specified register)
  - Debug support: INFO, ASSERT
    - INFO supports diagnostic generation
    - ASSERT generates error message during assembly if assertion is false


# Assembler Directives

- Examples:
  - Array DCD 1, 2, 3 ; array, not only 1 word
  - FailingGrade DCB "D, C-, C, or C+", 0
  - INFO 0, "Version 1.0" ; reserves 10 bytes
  - PerfectNumber EQU 10 ; Pythagora ~500 B.C
  - LF EQU 10 ; linefeed in ASCII
  - Area RN R5 ; R5 holds var Area

# Cortex ISA and Assembler: Hints

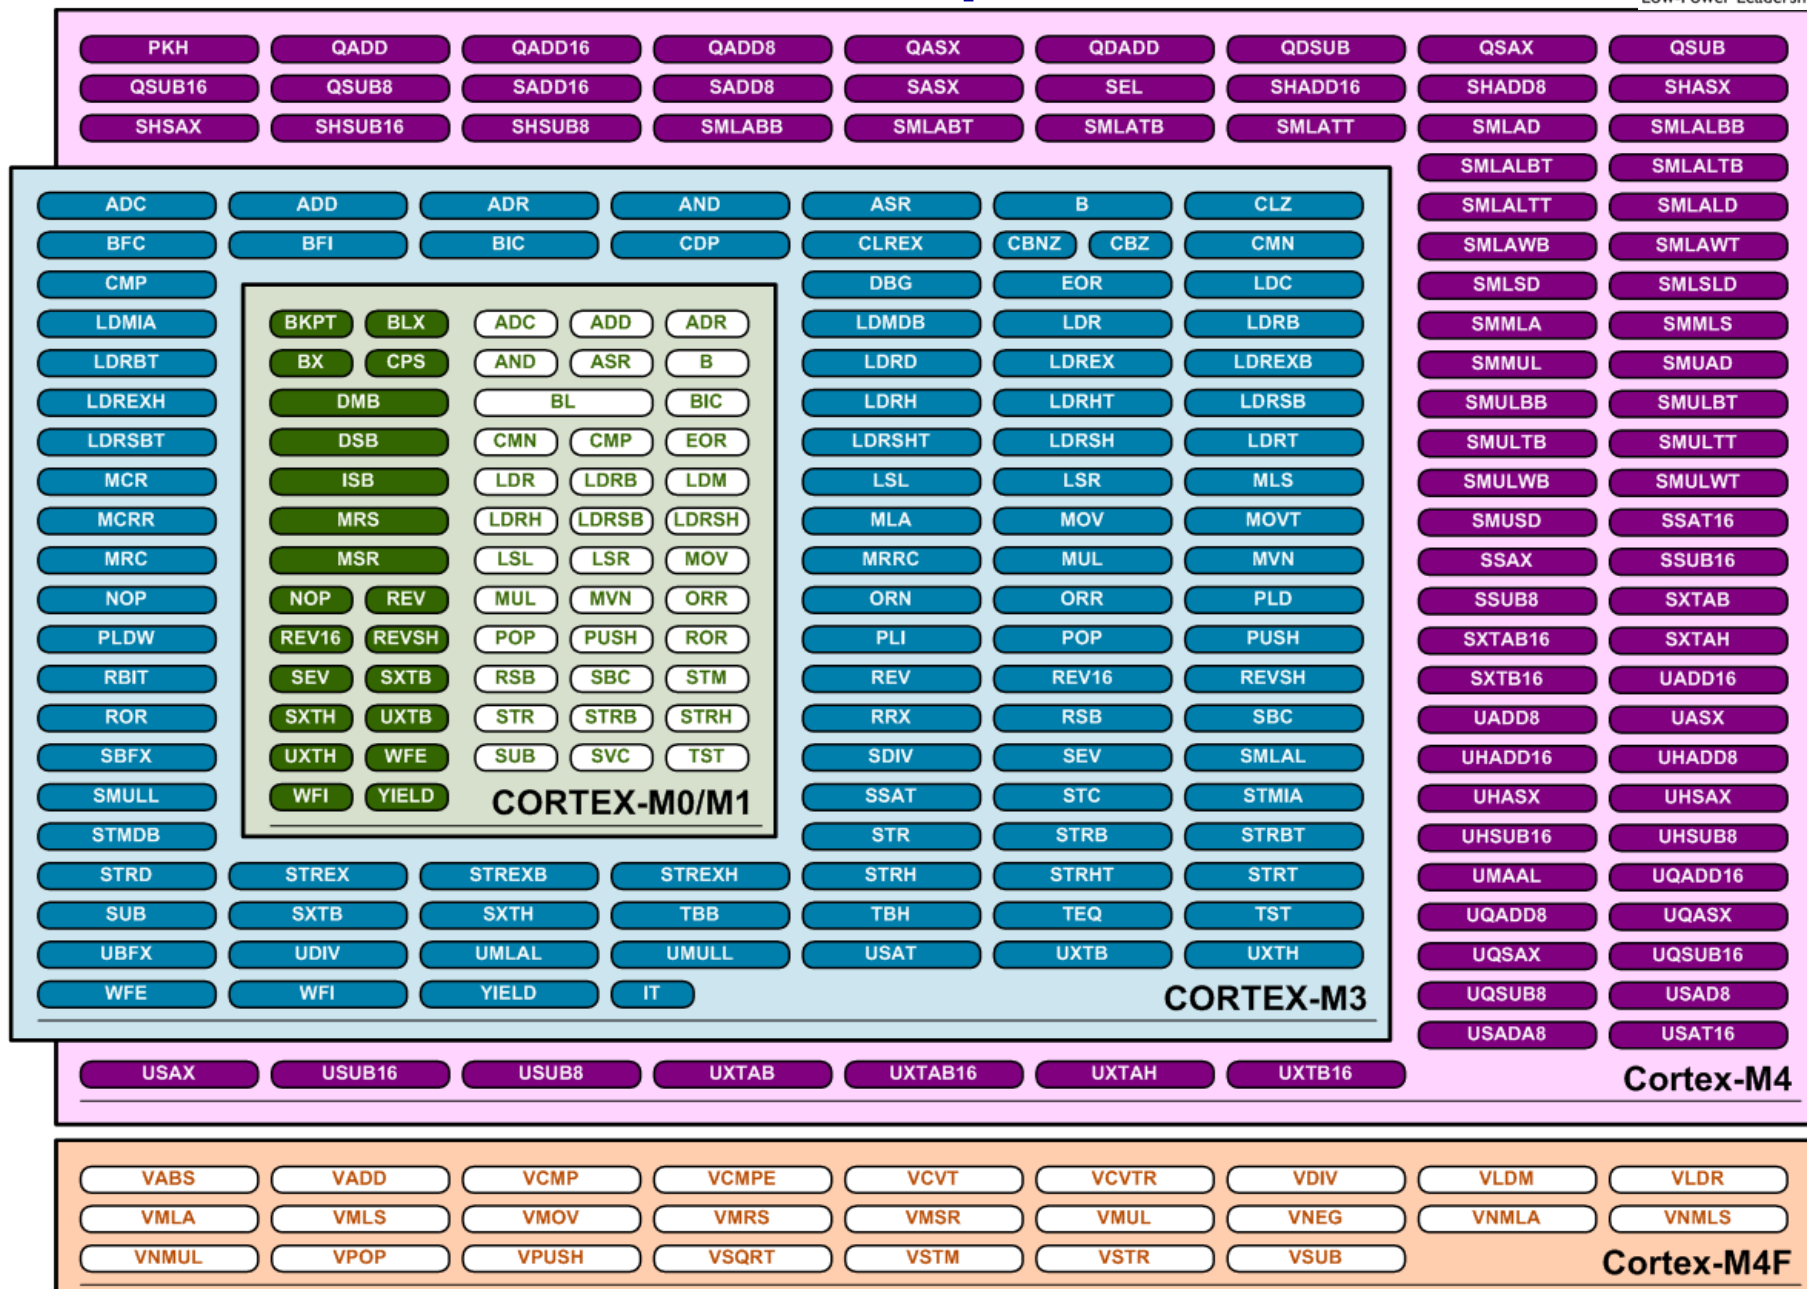
- Use conditional execution -> avoid branches
  - Simpler code
  - Full pipeline
- Useful instructions for bit-field operations:
  - BIC, BFC - sets some bits to zero
  - BFI - inserts bit fields to a word
  - AND, ORR, EOR - logical bit manipulations

# Cortex M3 ISA at Glance





# Cortex-M4 processors



# Embedded C

- C preprocessor
  - #define, #ifndef, #if, #ifdef, #else ...etc.
    - #define specifies flags for conditional compilation
    - All remaining preprocessor statements initiate conditional compilation
      - Example: #ifdef compiles a block of code if some condition is defined in the #define statement
  - #ifndef is used to prevent multiple includes
    - Use #ifndef if you want to include a new definition
  - Widely used in firmware (embedded SW) drivers

# C Preprocessor Examples

## Inline macro functions:

```
#define MIN(n,m) (((n) < (m)) ? (n) : (m))
```

```
#define MAX(n,m) (((n) < (m)) ? (m) : (n))
```

```
#define ABS(n) ((n < 0) ? -(n) : (n))
```

## Macro used to set LCD control

(## is used to actual arguments during macro expansion)

```
#define SET_VAL(x) LCD_Settings.P##x
```

## Nested macro definitions

```
#define SET(x, val) SET_VAL(x) = val
```

```
#define DEF_SET(x) SET(x, DS_P##x)
```

# Global variables

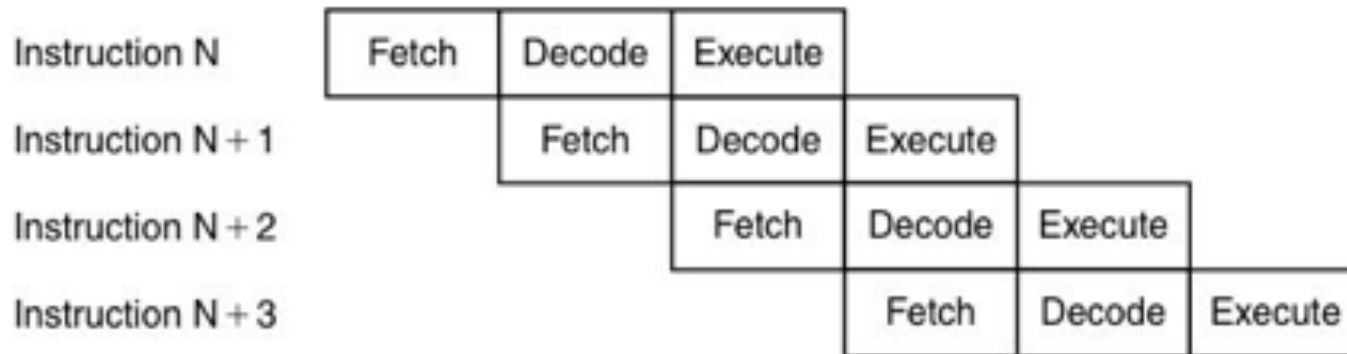
- Distinguish global variables from local by choosing appropriate naming convention
  - Example: `RX_Buffer_Gbl`
  - Stick to your convention throughout the program
- Use them as Software flags
  - Example: `PACKET_RECEIVED` – use capitals
- Have them all in ONE place
- Global variables are easy to observe during debug (“watch variables”)

# Cortex Pipeline

- CPU can execute most instructions in a single cycle
- Three-stage pipeline to increase flow rate of instructions
- Several operations to be performed simultaneously rather than serially
  - While the first instruction being executed,
  - the next instruction is being decoded
  - and the third is being fetched from memory

# Cortex Pipeline

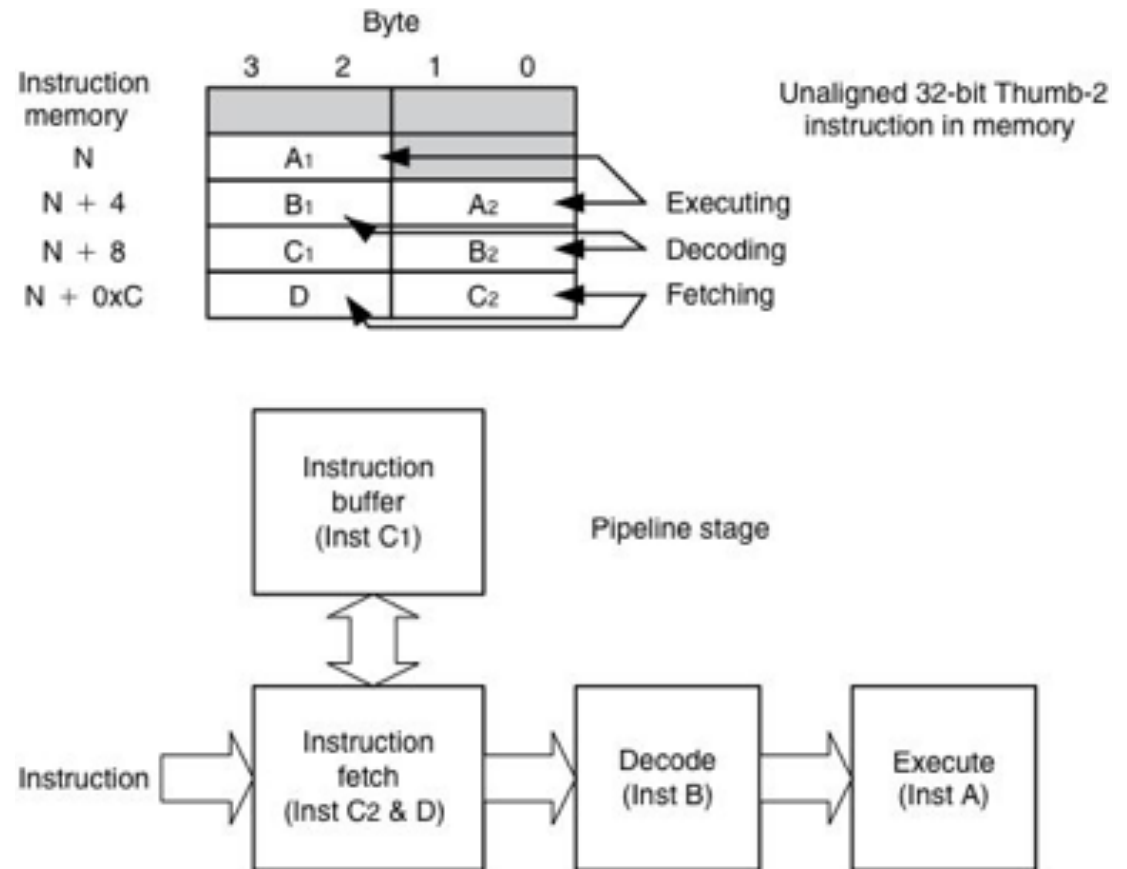
- PC points to the instruction being fetched, not executed
  - Debug tools will hide that from you
  - This is now part of the ARM Architecture and applies to all processors



# Cortex M Pipelining

- Three stages:

- Fetch
  - Instruction is fetched
  - Data is returned from instruction memory
- Decode
  - Generation of address using forward registers, branch forwarding
- Execution
  - Instruction is executed in a single pipeline stage



# Optimal Pipelining

- All operations are on registers (single cycle execution)
- Example: 6 clock cycles to execute 6 instructions
- Clock cycle per instruction (CPI = 1)

Cycle		1	2	3	4	5	6	7	8	9
Operation										
ADD		F	D	E						
SUB			F	D	E					
ORR				F	D	E				
AND					F	D	E			
ORR						F	D	E		
EOR							F	D	E	

F - Fetch   D - Decode   E - Execute