

ECSE 426 – Fall 2015

Microprocessor Systems

Dept. Electrical and Computer Engineering
McGill University



McGill

Outline

- Interrupt prioritization
- Exception Handling
- OS Support
 - Supervisor Calls
 - Pending Service Call
 - Exclusive Access

Interrupt Prioritization

- Each interrupt source has an 8-bit interrupt priority value
- Bits divided into pre-empting and non-pre-empting "sub-priority" levels
 - Sub-priority levels only apply if the pre-empting priority levels are the same
 - Software programmable PRIGROUP register field of the NVIC
 - chooses how many of the 8-bits are used for "group-priority" and how many are used for "sub-priority"
 - Group priority is the pre-empting priority
- In the STM32F10x 16 levels (4-bit) of priority are implemented:

PRIGROUP (3 Bits)	Binary Point (group.sub)		Preempting Priority (Group Priority)		Sub-Priority	
			Bits	Levels	Bits	Levels
011	4.0	gggg	4	16	0	0
100	3.1	gggs	3	8	1	2
101	2.2	ggss	2	4	2	4
110	1.3	gsss	1	2	3	8
111	0.4	ssss	0	0	4	16

Interrupt Prioritization

- Lower numbers are higher priority
- Hardware interrupt number is the lowest level of prioritization
 - IRQ3 is higher priority than IRQ4 if the priority registers are programmed the same
- In STM32F10x 16 levels (4-bit) of priority implemented:

PRIGROUP (3 Bits)	Binary Point (group.sub)		Preempting Priority (Group Priority)		Sub-Priority	
			Bits	Levels	Bits	Levels
011	4.0	gggg	4	16	0	0
100	3.1	gggs	3	8	1	2
101	2.2	ggss	2	4	2	4
110	1.3	gsss	1	2	3	8
111	0.4	ssss	0	0	4	16

Exception handlers in C

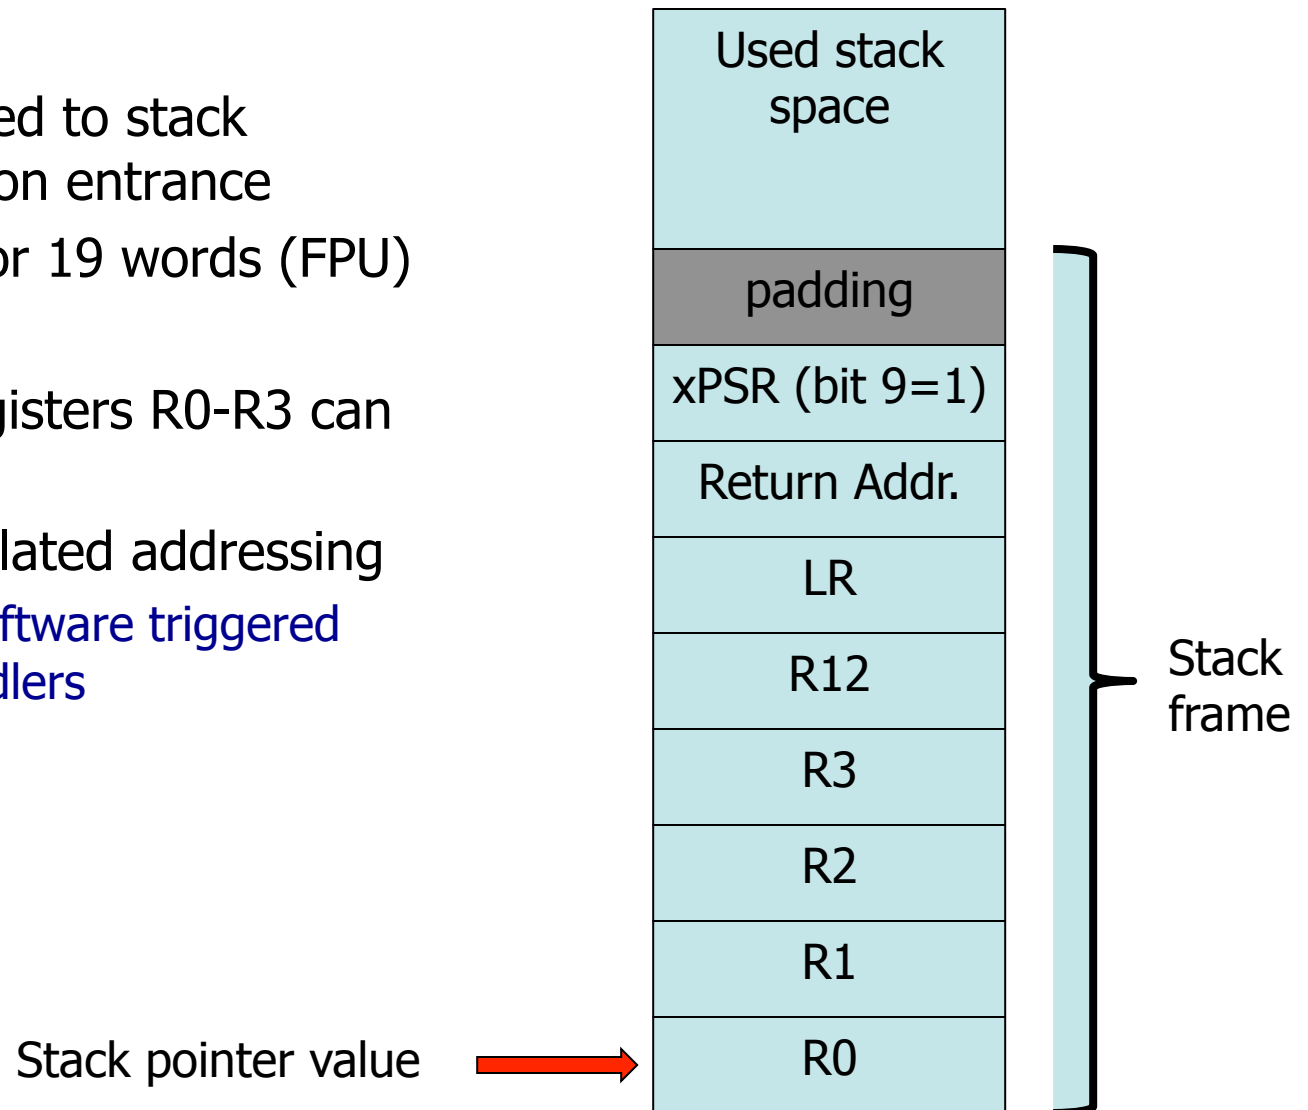
- C compilers follow AAPCS (ARM Architecture Procedure Call Standard)
- C function can modify R0 to R3, R12, R14 (LSR) and PSR
 - **caller saved** registers
 - Code that calls a subroutine must save them to memory if it needs them after the function call
- If C function needs to use R4-R11, it should save these registers onto stack memory and restore them before exiting
 - **Callee saved** registers
- Typically R0-R3 are input parameters, and R0 is return result (+ R1 for 64 bit result)

Exception handlers in C

- Exception mechanism needs to save R0-R3, R12, LR and PSR at exception entrance automatically
- Restore them at exception exit under control of processor's hardware
- Value of return address (PC) is not stored in LR as in normal C function calls
 - Exception mechanism puts an **EXC_RETURN** code in LR at exception entry
 - Value of return sequence also needs to be saved
- 8 registers need to be saved (+ floating point)

Stack Frame

- Block of data pushed to stack memory at exception entrance
- 8 words (no FPU) or 19 words (FPU)
- General purpose registers R0-R3 can be easily accessed
- Use stack pointer related addressing
- Pass information to software triggered interrupts or SVC handlers



EXC_RETURN

- When processor enters exception handler or interrupt service routine
 - Value of link register (LR) updated to a code called EXC_RETURN
 - Used to trigger exception return mechanism when loaded into program counter (PC)
- Some bits provide extra information about the exception sequence
 - 4: Stack frame type = 1 (8 words) or 0 (26 words)
 - 3: Return mode = 1 (Thread) or 0 (Handler)
 - 2: Return stack = 1 (Process stack) or 0 (Main stack)

Exception sequence

- Exception occurs
 - need to push registers into stack (form stack frame)
 - perform vector fetch
 - and start exception handler instruction fetch
- Multiple bus interfaces
 - reduce interrupt latency: do stacking and flash memory access in parallel
 - Stacking: data bus
 - Vector fetch + instruction fetch: instruction bus

Cortex Interrupts

Entry

Processor state automatically saved to stack over data bus (SYSTEM)

{R0-R3, R12, LR, PC, xPSR}

In parallel, ISR pre-fetched on the instruction bus (ICODE)

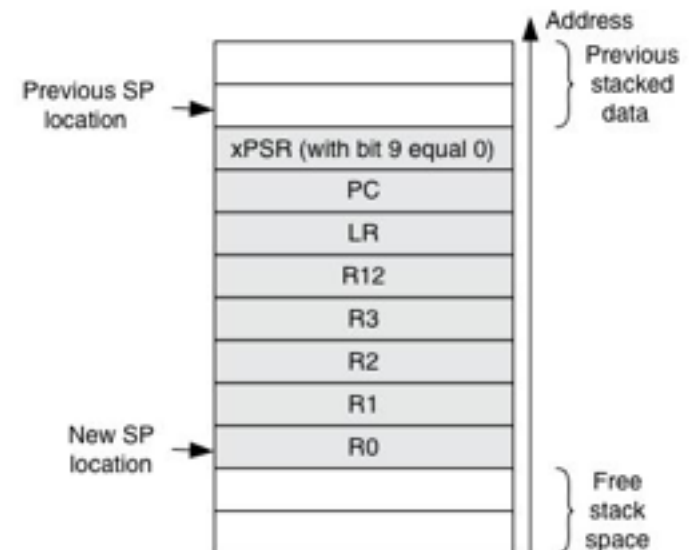
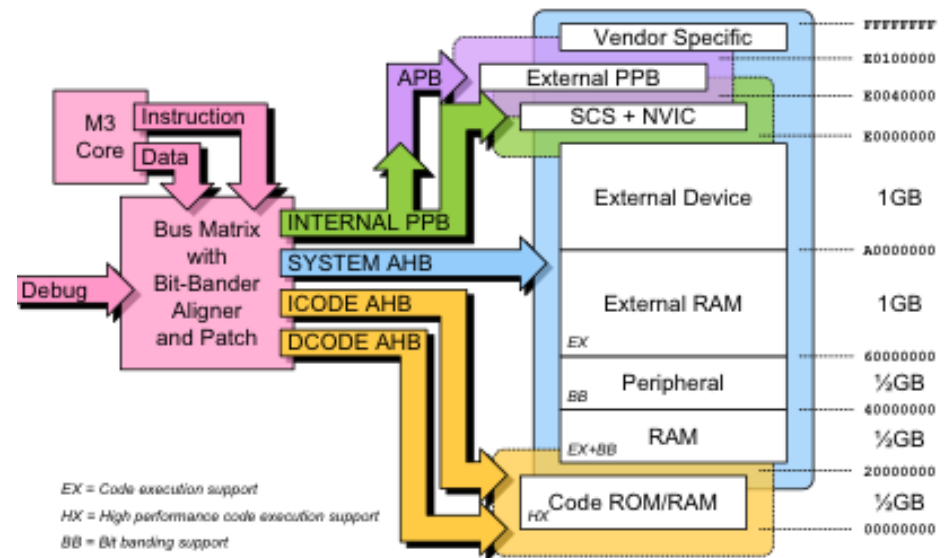
ISR ready to start executing as soon as stack PUSH complete

Exit

Processor state is automatically restored from the stack

In parallel, interrupted instruction is prefetched for execution upon POP

Stack POP can be interrupted -> new ISR immediately executed



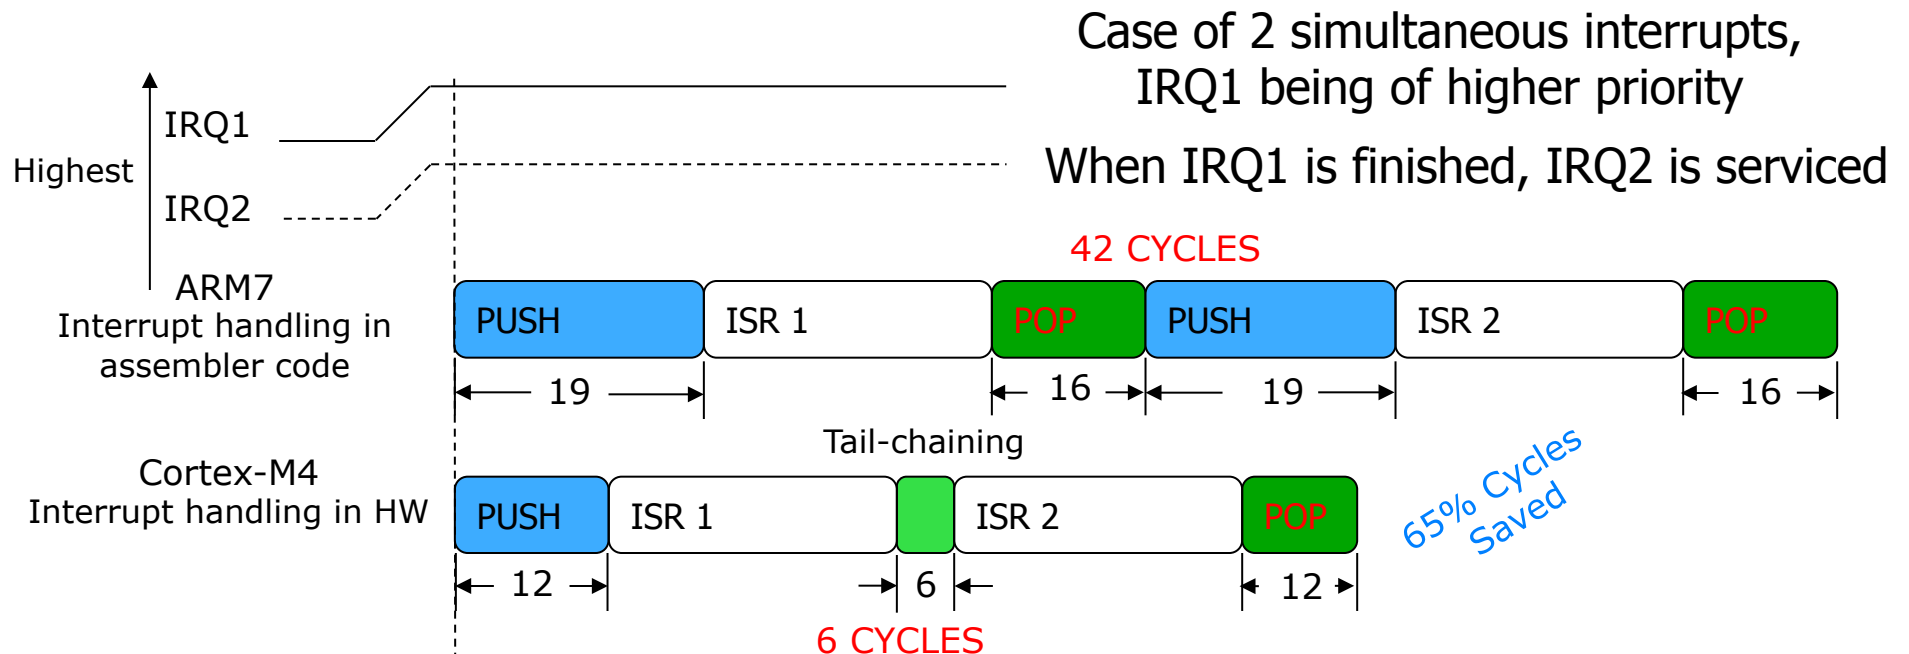
Interrupt latency

- Delay from start of interrupt request to start of interrupt handler execution
- If memory system has zero latency and bus system design allows vector fetch and stacking at same time, latency is only 12 clock cycles for Cortex-M
- Latency can be higher due to wait states in memory system
 - If processor is carrying out a memory transfer (including buffered write operations), transfer must be completed before exception sequence starts
- Cortex-M strives to minimize latency
 - Many operations handled by hardware, e.g. nested interrupt handling
 - No need to use software code to determine which interrupt to service
 - No need to use code to locate starting addresses of ISRs

Tail chaining

- When exception occurs, but processor is handling another exception of same or higher priority, exception enters pending state
- Processor finishes executing current exception handler, proceeds to process pending request
- Processor skips the steps of restoring registers from stack (unstacking) then pushing them back onto stack
- Timing gap between the two exceptions is considerably reduced
 - only 6 cycles for memory system with no-wait state
- Tail chaining optimization makes system more energy efficient
 - reduces number of stack accesses and hence amount of memory transfer

Fast Interrupts: Tail Chaining



In Cortex-M, ISR2 has only a 6-cycle delay. ISR2 has been **tail-chained**

ICODE (vector) and SYSTEM (stack) used in parallel

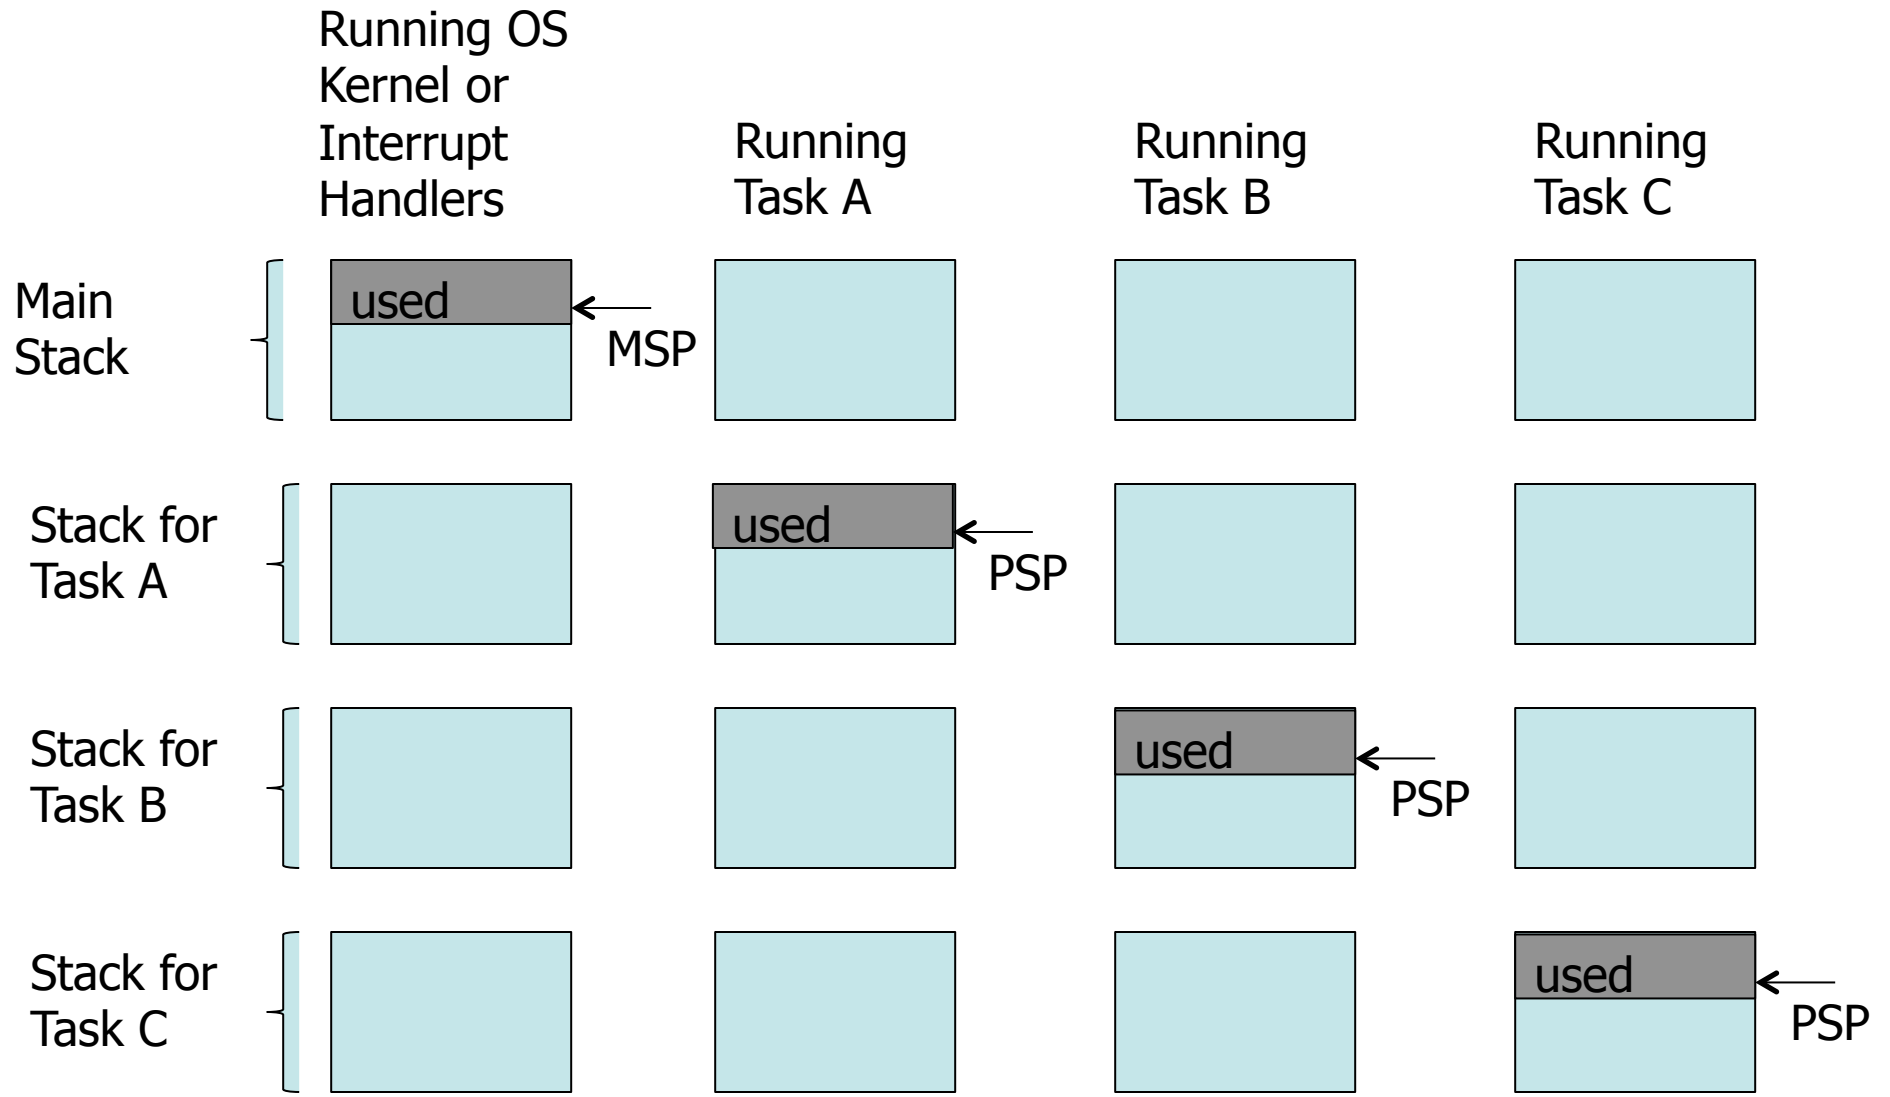
OS Support

- Shadowed stack pointer
 - Two stack pointers available
 - MSP used for OS kernel and interrupt handlers. PSP used by application tasks
- SysTick timer
 - Simple timer included inside processor.
 - Embedded OS can be used on a wide range of Cortex-M microcontrollers
- Supervisor Call (SVC) and Pendable Service Call (PendSV)
 - Allow context switching and other essential embedded OS operations
- Unprivileged execution level
 - Basic security model that restricts access rights of some application tasks
- Exclusive access
 - Exclusive load and store instructions are useful for semaphore and mutual exclusive (MUTEX) operations in the OS

Shadowed Stack Pointer

- Main Stack Pointer (MSP)
 - Default stack pointer
 - Used in Thread mode when CONTROL bit[1] (SPSEL) is 0
 - Always used in Handler mode
- Processor Stack Pointer (PSP)
 - Used in Thread mode when CONTROL bit[1] (SPSEL) is 1
- In systems with an embedded OS or RTOS
 - Exception handlers (including OS kernel) use MSP
 - Application tasks use PSP
 - Each application task has its own stack space
 - Context switching in OS updates PSP each time context is switched

Shadowed Stack Pointer



Shadowed Stack Pointer

- If application task has a problem that leads to stack corruption, stack of OS and other tasks likely remains intact
- Stack space for each task only needs to cover maximum stack usage plus one level of stack frame
- OS can use Memory Protection Unit (MPU) to define the stack region which an application task can use
 - If task has a stack overflow problem, MPU can trigger a MemManage fault exception
 - Prevent task from overwriting memory regions outside its allocated stack space

Context Switching

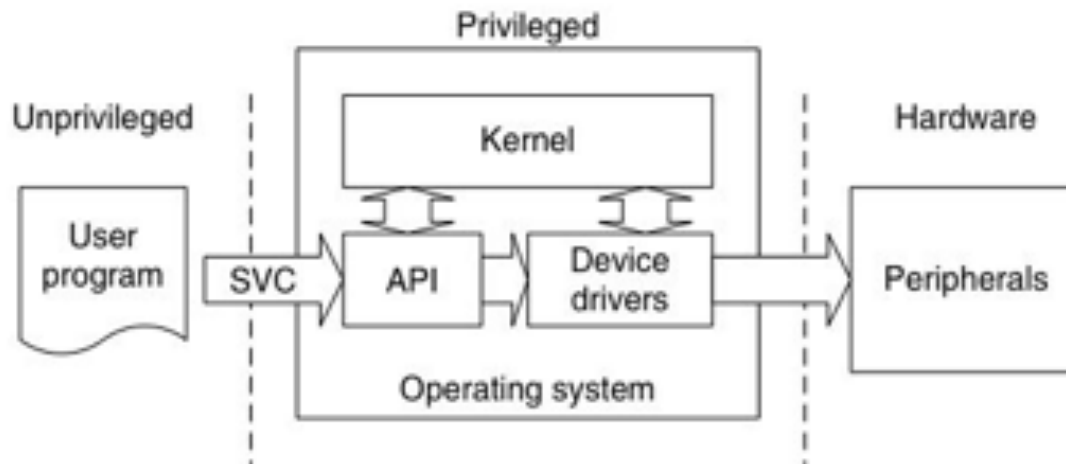
- OS needs to switch between different tasks
- Context switching usually carried out in the PendSV exception handler
- Can be triggered by a periodic SysTick exception
- Inside context switching operation
 - Save current status of registers in the current task
 - Save the current PSP value
 - Set the PSP value to the last SP value for the next task
 - Restore the last values for the next task
 - Use exception return to switch to the task
- Context switching is carried out in PendSV
 - Typically at lowest priority level
 - Prevents context switching from happening in the middle of an interrupt handler

Supervisor Calls: SVC & PendSV

- SVC (Supervisor Call) and PendSV (Pendable Service Call)
 - Important for OS designs
- SVC Instruction
 - Keil/ARM: `__svc`
 - Portable; hardware abstraction
 - Can write to NVIC using a software trigger interrupt register, but several instructions might execute while the interrupt is pending
 - With `__svc`, the SVC handler executes immediately (except when another higher priority exception arrives)
- Can be used as API to allow application tasks to access system resources

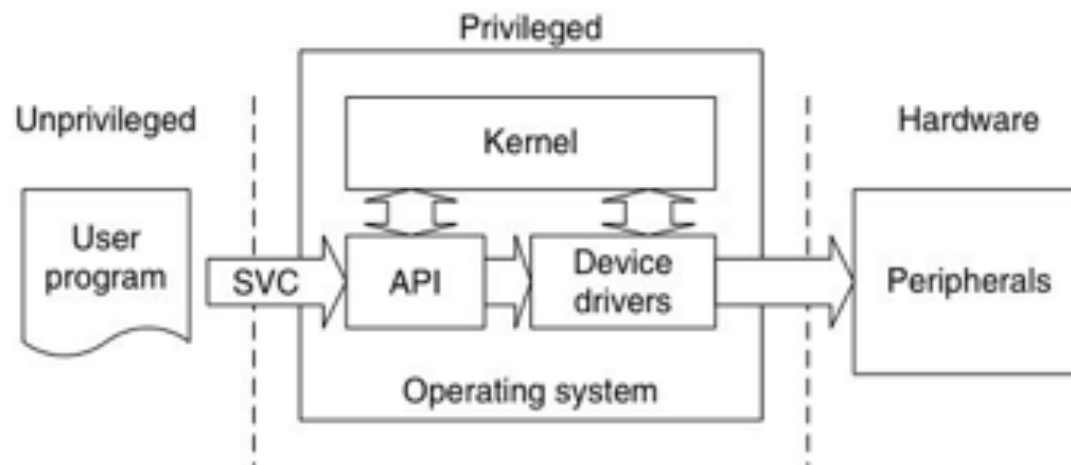
Supervisor Calls: SVC

- Systems with high reliability requirements
 - Some hardware set up to be privileged access only
 - Embedded system is more robust and secure



Supervisor Calls: SVC

- Application tasks
 - run in unprivileged access level
 - can only access protected hardware resources via services from OS
 - Cannot gain unauthorized access to critical hardware
 - Do not need to know programming details of underlying hardware
 - Can be developed independently of OS (don't need to know exact address of OS service functions – only the SVC service number and parameters)
 - Hardware level programming handled by device drivers



SVC Handling

- SVC instruction (`__svc`)
 - Needs immediate value (`SVC #0x3; Call SVC function 3`)
 - SVC exception handler extracts parameter and knows which action to perform
- Procedure
 - Determine if the calling program was using main stack or process stack?
 - Read the program counter value from appropriate stack
 - Read instruction from that address (masking out unnecessary bits)

`SVC_Handler`

```
TST lr, #4                ; Test bit 2 of EXC_RETURN
ITE EQ
MRSNE R0, MSP             ; if 0, stacking using MSP, copy to R0
MRSEQ R0, PSP             ; if 1, stacking using PSP, copy to R0
LDR R0, [R0, #24]         ; get stacked PC from stack frame
                          ; stacked PC = address of instruction after SVC)
LDRB R0, [R0, #-2]        ; get first byte of the SVC instruction
                          ; now SVC number is in R0
```

SVC Handling

- In C, we break it into two parts
- Can't check the value of LR (EXC_RETURN) in C
- Use assembly inline (`__asm`)

```
__asm void SVC_Handler(void)
{
    TST LR, #4           ; Test bit 2 of EXC_RETURN
    ITE EQ
    MRSNE R0, MSP        ; if 0, stacking using MSP, copy to R0
    MRSEQ R0, PSP        ; if 1, stacking using PSP, copy to R0
    B _cpp(SVC_Handler_C)
    ALIGN 4
}
```

SVC Handling

```
void SVC_Handler_C(unsigned int * svc_args)
{
    uint8_t svc_number;
    uint32_t stacked_r0, stacked_r1, stacked_r2, stacked_r3;

    svc_number = ((char *) svc_args[6])[-2];
    // Memory[(Stacked PC)-2]
    stacked_r0 = svc_args[0];
    stacked_r1 = svc_args[1];
    stacked_r2 = svc_args[2];
    stacked_r3 = svc_args[3];

    // other processing
    ...
    // Return result (e.g. sum of first two elements)
    svc_args[0] = stacked_r0 + stacked_r1;
    return;
}
```

Source : J. Yiu, "The Definitive Guide to ARM Cortex-M3 and Cortex-M4 Processors"

SVC Handling

- Passing the address of the stack frame allows the C handler to extract any information it needs
- Essential if you want to pass parameters to an SVC service and get a return value
- A higher priority interrupt could be executed first and change the values of R0, R1, etc.
- Using the stack frame ensures your SVC handler gets the correct input parameters

PendSV

- Pended Service Call (exception type 14)
 - Programmable priority level
 - Triggered by writing to Interrupt Control and Status Register (ICSR)
 - Unlike SVC, it is not precise. Pending status can be set in a higher priority exception handler
- Execution of OS kernel (and context switching) triggered by
 - SVC call from an application task (e.g. task is stalled because it is waiting for data or an event)
 - Periodic SysTick exception
- PendSV is lowest priority exception
 - Context switching delayed until all other IRQ handlers have finished
 - OS can set pending status of PendSV & carry out context switching in PendSV exception handler

Exclusive Access

- Multi-tasking system: tasks need to share limited resources
 - For example, one console output
- Tasks need to “lock” a resource and then “free” it after use
 - Usually based on software variables
 - If lock variable is set, other tasks can see that the resource is locked
- Lock variable is called a **semaphore**
 - If only one resource is available, also called **Mutual Exclusive (MUTEX)**
 - In general, semaphores can support multiple tokens
 - e.g. one for each channel of a communication stack
 - Semaphore implemented as a token counter
 - Each task decrements the semaphore when it needs the resource

Exclusive Access (2)

- Decrement of counter is not atomic
 - One instruction to read variable
 - One instruction to decrement it
 - One instruction to write back to memory
- Context switching may occur between read and write
 - Simplest approach: disable context switching when handling semaphores
 - Can increase latency and only works for single processor designs
 - Multi-processor: tasks on different processors can try to decrement semaphore variable at same time

Exclusive Access: Local Monitor

- Cortex-M4 supports feature called **exclusive access**
 - Semaphores are read and written using exclusive load and exclusive store
 - If during store, access cannot be guaranteed to be exclusive, exclusive store fails
- Processor has small hardware unit called the **local monitor**
 - Normally in **Open Access** state
 - Exclusive load: switches to **Exclusive Access** state
 - Exclusive store: only executes if local monitor in Exclusive Access state
 - STREX fails if
 - CLREX has been executed (switching local monitor back to Open)
 - Context switch has occurred (interrupt)
 - No LDREX was executed earlier
 - External hardware returns an exclusive fail status
 - Multiprocessor systems need a global exclusive access monitor

Summary

- Exception handling
 - Exception return code is inserted into link register
 - Provides key information about the exception
 - Handling is built around the stack frame
- Cortex M4 and the Discovery board provide multiple resources to support operating system functionality
 - Shadowed stack: different stacks for different tasks
 - Supervisor calls: OS provides hardware access services
 - Pending service call: context switching
 - Exclusive access: local monitor and exclusive load/store